

▼ Data preprocessing

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Steps of Data Preprocessing

- Basic preprocessing:- Removing commas, drop columns like indexes, identifiers, name, address etc.
- Handling missing values
- Encoding of catogorical columns using dummy encoding and frequency encoding based on cardinality
- Scaling of the features using StandardScaler on both numerical and transformed columns

▼ 1- Basic preprocessing like drop columns

```
df=pd.read_csv('Healthcare Providers.csv')
df.head()
```

↗

	index	National Provider Identifier	Last Name/Organization Name of the Provider	First Name of the Provider	Middle Initial of the Provider	Credentials of the Provider	Gende of th Provide
0	8774979	1891106191	UPADHYAYULA	SATYASREE	NaN	M.D.	
1	3354385	1346202256	JONES	WENDY	P	M.D.	
2	3001884	1306820956	DUROCHER	RICHARD	W	DPM	I
3	7594822	1770523540	FULLARD	JASPER	NaN	MD	I
4	746159	1073627758	PERROTTI	ANTHONY	E	DO	I

5 rows × 27 columns

```
df.info()
```

↗

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   index                                100000 non-null  int64
1   National Provider Identifier          100000 non-null  int64
2   Last Name/Organization Name of the Provider  100000 non-null  object
3   First Name of the Provider           95745 non-null   object
4   Middle Initial of the Provider        70669 non-null   object
5   Credentials of the Provider           92791 non-null   object
6   Gender of the Provider                95746 non-null   object
7   Entity Type of the Provider           100000 non-null  object
8   Street Address 1 of the Provider       100000 non-null  object
9   Street Address 2 of the Provider       40637 non-null   object
10  City of the Provider                  100000 non-null  object
```

```

11 Zip Code of the Provider          100000 non-null float64
12 State Code of the Provider        100000 non-null object
13 Country Code of the Provider      100000 non-null object
14 Provider Type                     100000 non-null object
15 Medicare Participation Indicator   100000 non-null object
16 Place of Service                  100000 non-null object
17 HCPCS Code                        100000 non-null object
18 HCPCS Description                  100000 non-null object
19 HCPCS Drug Indicator               100000 non-null object
20 Number of Services                100000 non-null object
21 Number of Medicare Beneficiaries  100000 non-null object
22 Number of Distinct Medicare Beneficiary/Per Day Services 100000 non-null object
23 Average Medicare Allowed Amount   100000 non-null object
24 Average Submitted Charge Amount   100000 non-null object
25 Average Medicare Payment Amount   100000 non-null object
26 Average Medicare Standardized Amount 100000 non-null object
dtypes: float64(1), int64(2), object(24)
memory usage: 20.6+ MB

```

## df.columns

```

Index(['index', 'National Provider Identifier',
      'Last Name/Organization Name of the Provider',
      'First Name of the Provider', 'Middle Initial of the Provider',
      'Credentials of the Provider', 'Gender of the Provider',
      'Entity Type of the Provider', 'Street Address 1 of the Provider',
      'Street Address 2 of the Provider', 'City of the Provider',
      'Zip Code of the Provider', 'State Code of the Provider',
      'Country Code of the Provider', 'Provider Type',
      'Medicare Participation Indicator', 'Place of Service', 'HCPCS Code',
      'HCPCS Description', 'HCPCS Drug Indicator', 'Number of Services',
      'Number of Medicare Beneficiaries',
      'Number of Distinct Medicare Beneficiary/Per Day Services',
      'Average Medicare Allowed Amount', 'Average Submitted Charge Amount',
      'Average Medicare Payment Amount',
      'Average Medicare Standardized Amount'],
      dtype='object')

```

```

columns=['index', 'National Provider Identifier',
        'Last Name/Organization Name of the Provider',
        'First Name of the Provider', 'Middle Initial of the Provider', 'Street Address 1 of the Provider',
        'Street Address 2 of the Provider', 'Zip Code of the Provider', 'HCPCS Description']

```

```
df=df.drop(columns,axis=1)
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 18 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Credentials of the Provider                92791 non-null  object
1   Gender of the Provider                    95746 non-null  object
2   Entity Type of the Provider               100000 non-null object
3   City of the Provider                     100000 non-null object
4   State Code of the Provider                100000 non-null object
5   Country Code of the Provider              100000 non-null object
6   Provider Type                           100000 non-null object
7   Medicare Participation Indicator           100000 non-null object
8   Place of Service                         100000 non-null object
9   HCPCS Code                              100000 non-null object
10  HCPCS Drug Indicator                     100000 non-null object
11  Number of Services                       100000 non-null object
12  Number of Medicare Beneficiaries          100000 non-null object
13  Number of Distinct Medicare Beneficiary/Per Day Services 100000 non-null object
14  Average Medicare Allowed Amount           100000 non-null object
15  Average Submitted Charge Amount           100000 non-null object
16  Average Medicare Payment Amount           100000 non-null object
17  Average Medicare Standardized Amount      100000 non-null object
dtypes: object(18)
memory usage: 13.7+ MB

```

```
# converting the values like M.D. as MD As it has the same meaning
```

```
df['Credentials of the Provider']=df['Credentials of the Provider'].str.replace(".", "")
```

```
# Preprocessing for numerical columns
df.iloc[:,11:]=df.iloc[:,11:].apply(lambda x: x.str.replace(',', ''))
```

```
# Converting numerical columns of 'object' datatypes as 'Float'
object_cols = df.iloc[:,11:].columns
df[object_cols] = df[object_cols].apply(lambda x: x.astype(float))
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 18 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Credentials of the Provider           92791 non-null  object
 1   Gender of the Provider                95746 non-null  object
 2   Entity Type of the Provider           100000 non-null object
 3   City of the Provider                  100000 non-null object
 4   State Code of the Provider            100000 non-null object
 5   Country Code of the Provider          100000 non-null object
 6   Provider Type                        100000 non-null object
 7   Medicare Participation Indicator       100000 non-null object
 8   Place of Service                     100000 non-null object
 9   HCPCS Code                           100000 non-null object
10   HCPCS Drug Indicator                  100000 non-null object
11   Number of Services                   100000 non-null float64
12   Number of Medicare Beneficiaries      100000 non-null float64
13   Number of Distinct Medicare Beneficiary/Per Day Services 100000 non-null float64
14   Average Medicare Allowed Amount       100000 non-null float64
15   Average Submitted Charge Amount       100000 non-null float64
16   Average Medicare Payment Amount       100000 non-null float64
17   Average Medicare Standardized Amount 100000 non-null float64
dtypes: float64(7), object(11)
memory usage: 13.7+ MB
```

## ✓ Handling missing values

```
df.isnull().sum()
```

```
Credentials of the Provider           7209
Gender of the Provider                 4254
Entity Type of the Provider              0
City of the Provider                    0
State Code of the Provider               0
Country Code of the Provider             0
Provider Type                           0
Medicare Participation Indicator          0
Place of Service                        0
HCPCS Code                              0
HCPCS Drug Indicator                    0
Number of Services                      0
Number of Medicare Beneficiaries         0
Number of Distinct Medicare Beneficiary/Per Day Services 0
Average Medicare Allowed Amount          0
Average Submitted Charge Amount          0
Average Medicare Payment Amount          0
Average Medicare Standardized Amount     0
dtype: int64
```

```
# fill the missing the values with mode of the columns for 'credential of the provider'
df["Credentials of the Provider"] = df["Credentials of the Provider"].fillna(df["Credentials of th
# fill the missing values of 'Gender of the provider' by creating another category for the organiz
df["Gender of the Provider"] = df["Gender of the Provider"].fillna('0')
```

```

columns=['Credentials of the Provider', 'Gender of the Provider',
        'Entity Type of the Provider', 'City of the Provider',
        'State Code of the Provider', 'Country Code of the Provider',
        'Provider Type', 'Medicare Participation Indicator', 'Place of Service',
        'HCPCS Code', 'HCPCS Drug Indicator']
for i in columns:
    if df[i].nunique() > 5:
        print(f"categorical values in {i} :",df[i].nunique())
print("\n")
for i in columns:
    if df[i].nunique() < 5:
        print(f"categorical values in {i} :",df[i].nunique())

```

```

categorical values in Credentials of the Provider : 1539
categorical values in City of the Provider : 5846
categorical values in State Code of the Provider : 58
categorical values in Provider Type : 90
categorical values in HCPCS Code : 2631

categorical values in Gender of the Provider : 3
categorical values in Entity Type of the Provider : 2
categorical values in Country Code of the Provider : 4
categorical values in Medicare Participation Indicator : 2
categorical values in Place of Service : 2
categorical values in HCPCS Drug Indicator : 2

```

df.columns

```

Index(['Credentials of the Provider', 'Gender of the Provider',
      'Entity Type of the Provider', 'City of the Provider',
      'State Code of the Provider', 'Country Code of the Provider',
      'Provider Type', 'Medicare Participation Indicator', 'Place of Service',
      'HCPCS Code', 'HCPCS Drug Indicator', 'Number of Services',
      'Number of Medicare Beneficiaries',
      'Number of Distinct Medicare Beneficiary/Per Day Services',
      'Average Medicare Allowed Amount', 'Average Submitted Charge Amount',
      'Average Medicare Payment Amount',
      'Average Medicare Standardized Amount'],
      dtype='object')

```

### ✓ 3- Encoding of categorical columns

```

# Columns to be frequency encoded
frequency_encode_cols = ['Credentials of the Provider', 'City of the Provider',
                        'State Code of the Provider', 'Provider Type', 'HCPCS Code']

# Function to perform frequency encoding
def frequency_encoding(df, columns):
    for column in columns:
        frequency = df[column].value_counts()
        df[column] = df[column].map(frequency)
    return df

# Apply frequency encoding
encoded_data = frequency_encoding(df, frequency_encode_cols)
encoded_data.head()

```



	Credentials of the Provider	Gender of the Provider	Entity Type of the Provider	City of the Provider	State Code of the Provider	Country Code of the Provider	Provider Type	Medicare Participation Indicator
0	73827	F	I	500	1997	US	11366	
1	73827	F	I	209	3725	US	1028	
2	1915	M	I	10	1403	US	2027	
3	73827	M	I	317	1997	US	11366	
4	6176	M	I	51	7263	US	11366	

```
new_df=pd.get_dummies(encoded_data,drop_first=True)
new_df.head()
```



	Credentials of the Provider	City of the Provider	State Code of the Provider	Provider Type	HCPCS Code	Number of Services	Number of Medicare Beneficiaries	Number of Distinct Medicare Beneficiary/Per Day Services
0	73827	500	1997	11366	1297	27.0	24.0	
1	73827	209	3725	1028	243	175.0	175.0	
2	1915	10	1403	2027	44	32.0	13.0	
3	73827	317	1997	11366	460	20.0	18.0	
4	6176	51	7263	11366	732	33.0	24.0	

5 rows × 21 columns

```
new_df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 21 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Credentials of the Provider                                           100000 non-null  int64
1   City of the Provider                                                  100000 non-null  int64
2   State Code of the Provider                                           100000 non-null  int64
3   Provider Type                                                         100000 non-null  int64
4   HCPCS Code                                                            100000 non-null  int64
5   Number of Services                                                    100000 non-null  float64
6   Number of Medicare Beneficiaries                                     100000 non-null  float64
7   Number of Distinct Medicare Beneficiary/Per Day Services            100000 non-null  float64
8   Average Medicare Allowed Amount                                     100000 non-null  float64
9   Average Submitted Charge Amount                                     100000 non-null  float64
10  Average Medicare Payment Amount                                     100000 non-null  float64
11  Average Medicare Standardized Amount                               100000 non-null  float64
12  Gender of the Provider_M                                             100000 non-null  bool
13  Gender of the Provider_O                                             100000 non-null  bool
14  Entity Type of the Provider_O                                       100000 non-null  bool
15  Country Code of the Provider_JP                                     100000 non-null  bool
16  Country Code of the Provider_TR                                     100000 non-null  bool
17  Country Code of the Provider_US                                     100000 non-null  bool
18  Medicare Participation Indicator_Y                                   100000 non-null  bool
19  Place of Service_O                                                  100000 non-null  bool
20  HCPCS Drug Indicator_Y                                              100000 non-null  bool
dtypes: bool(9), float64(7), int64(5)
memory usage: 10.0 MB
```

3- Standardized the data

```
# Standardized the data
from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
scaled_data=ss.fit_transform(new_df)

final_df=pd.DataFrame(scaled_data,columns=final_df.columns)
```

#### ✓ 4- Pricipal component analysis and visualization of first two PCA components

```
# Performing PCA
from sklearn.decomposition import PCA
pca=PCA(n_components=None)
df_trf=pca.fit_transform(final_df)
```

df\_trf

```
array([[ 5.65151082e-01, -3.68919853e-01, -4.17790832e-01, ...,
         1.61154901e-02, -8.33555352e-03, -5.58553832e-16],
       [ 4.14901510e-01,  2.23484065e-01,  3.67014911e-01, ...,
         1.13743376e-01, -9.02062229e-02, -5.26233615e-15],
       [-3.09223405e-01, -2.33799502e-01, -8.54630895e-02, ...,
        -3.67317244e-02,  1.43151872e-02,  9.82668367e-16],
       ...,
       [-4.41501798e-01, -4.44611890e-01, -5.94041473e-01, ...,
        -2.91651451e-03, -2.43013546e-03, -3.08753651e-16],
       [-7.12466671e-01, -8.24807954e-02,  1.86182714e-01, ...,
         5.53457407e-03, -2.91659880e-02, -1.34879998e-15],
       [-4.90889805e-01,  4.23888886e-01,  3.56201987e-01, ...,
        -7.86910122e-03, -3.04036065e-03, -1.55310233e-16]])
```

```
# Create a DataFrame with the principal components
pca_df = pd.DataFrame(data=df_trf, columns=[f'PC{i}' for i in range(1, 22)])
```

```
# Plot the First Two Principal Components
plt.figure(figsize=(8, 6))
plt.scatter(pca_df['PC1'], pca_df['PC2'], alpha=0.5)
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('PCA - First Two Principal Components')
plt.grid()
plt.show()

pca_df.head()
```

