

## milestone-4

July 16, 2024

### 1 Preprocessing :-

Loading the dataset

```
[ ]: # prompt: load the dataset /content/Healthcare Providers.csv

import pandas as pd
df = pd.read_csv('/content/Healthcare Providers.csv')
```

Columns present in the dataset

```
[ ]: # prompt: display the columns present in the /content/Healthcare Providers.csv

import pandas as pd

# Loading the dataset
df = pd.read_csv('/content/Healthcare Providers.csv')

# Columns present in the dataset
print(df.columns)
```

```
Index(['index', 'National Provider Identifier',
      'Last Name/Organization Name of the Provider',
      'First Name of the Provider', 'Middle Initial of the Provider',
      'Credentials of the Provider', 'Gender of the Provider',
      'Entity Type of the Provider', 'Street Address 1 of the Provider',
      'Street Address 2 of the Provider', 'City of the Provider',
      'Zip Code of the Provider', 'State Code of the Provider',
      'Country Code of the Provider', 'Provider Type',
      'Medicare Participation Indicator', 'Place of Service', 'HCPCS Code',
      'HCPCS Description', 'HCPCS Drug Indicator', 'Number of Services',
      'Number of Medicare Beneficiaries',
      'Number of Distinct Medicare Beneficiary/Per Day Services',
      'Average Medicare Allowed Amount', 'Average Submitted Charge Amount',
      'Average Medicare Payment Amount',
      'Average Medicare Standardized Amount'],
      dtype='object')
```

Removing the columns which are not needed

```
[ ]: last_dataset = df.drop(['Street Address 2 of the Provider',
                            'Street Address 1 of the Provider',
                            'Zip Code of the Provider',
                            'index',
                            'National Provider Identifier',
                            'Last Name/Organization Name of the Provider',
                            'First Name of the Provider'], axis=1)

# save the dataset
last_dataset.to_csv('last_dataset.csv', index=False)
```

display unique values for columns

```
[ ]: for col in last_dataset.columns:
      print(col, ":", last_dataset[col].nunique())
```

```
Middle Initial of the Provider : 29
Credentials of the Provider : 1854
Gender of the Provider : 2
Entity Type of the Provider : 2
City of the Provider : 5846
State Code of the Provider : 58
Country Code of the Provider : 4
Provider Type : 90
Medicare Participation Indicator : 2
Place of Service : 2
HCPCS Code : 2631
HCPCS Description : 2455
HCPCS Drug Indicator : 2
Number of Services : 2748
Number of Medicare Beneficiaries : 1274
Number of Distinct Medicare Beneficiary/Per Day Services : 1979
Average Medicare Allowed Amount : 49629
Average Submitted Charge Amount : 38088
Average Medicare Payment Amount : 83367
Average Medicare Standardized Amount : 76237
```

performing onehot encoding

```
[ ]: import pandas as pd
      from sklearn.preprocessing import OneHotEncoder

      # Create the encoder.
      encoder = OneHotEncoder(handle_unknown="ignore")

      # Fit the encoder on the categorical features.
      encoder.fit(last_dataset[['Gender of the Provider',
                               'Entity Type of the Provider',
                               'Medicare Participation Indicator',
```

```

        'HCPCS Drug Indicator']])

# Transform the categorical features into one-hot encoded features.
encoded_features = encoder.transform(last_dataset[['Gender of the Provider',
        'Entity Type of the Provider',
        'Medicare Participation_
        ↪Indicator'],
        'HCPCS Drug Indicator']]).
        ↪toarray()

# Create column names for the one-hot encoded features.
encoded_feature_names = encoder.get_feature_names_out(['Gender of the Provider',
        'Entity Type of the_
        ↪Provider',
        'Medicare Participation_
        ↪Indicator'],
        'HCPCS Drug Indicator'])

# Create a new DataFrame with the one-hot encoded features.
encoded_df = pd.DataFrame(encoded_features, columns=encoded_feature_names)

# Concatenate the original DataFrame and the one-hot encoded DataFrame.
last_dataset = pd.concat([last_dataset, encoded_df], axis=1)

# Drop the original categorical features.
last_dataset = last_dataset.drop(['Gender of the Provider',
        'Entity Type of the Provider',
        'Medicare Participation Indicator',
        'HCPCS Drug Indicator'], axis=1)

# Save the updated dataset.
last_dataset.to_csv('/content/last_dataset.csv', index=False)

```

```
[ ]: last_dataset.head()
```

```
[ ]: Middle Initial of the Provider Credentials of the Provider \
0           NaN           M.D.
1           P           M.D.
2           W           DPM
3           NaN           MD
4           E           DO

City of the Provider State Code of the Provider \
0      SAINT LOUIS           MO
1    FAYETTEVILLE           NC
2    NORTH HAVEN           CT
3    KANSAS CITY           MO
```

4

JUPITER

FL

	Country Code of the Provider	Provider Type	Place of Service	\
0	US	Internal Medicine		F
1	US	Obstetrics & Gynecology		0
2	US	Podiatry		0
3	US	Internal Medicine		0
4	US	Internal Medicine		0

	HCPCS Code	HCPCS Description	\
0	99223	Initial hospital inpatient care, typically 70 ...	
1	G0202	Screening mammography, bilateral (2-view study...	
2	99348	Established patient home visit, typically 25 m...	
3	81002	Urinalysis, manual test	
4	96372	Injection beneath the skin or into muscle for ...	

	Number of Services	... Average Medicare Standardized Amount	\
0	27	...	160.90888889
1	175	...	135.31525714
2	32	...	60.5959375
3	20	...	3.43
4	33	...	19.057575758

	Gender of the Provider_F	Gender of the Provider_M	\
0	1.0	0.0	
1	1.0	0.0	
2	0.0	1.0	
3	0.0	1.0	
4	0.0	1.0	

	Gender of the Provider_nan	Entity Type of the Provider_I	\
0	0.0	1.0	
1	0.0	1.0	
2	0.0	1.0	
3	0.0	1.0	
4	0.0	1.0	

	Entity Type of the Provider_0	Medicare Participation Indicator_N	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	

	Medicare Participation Indicator_Y	HCPCS Drug Indicator_N	\
0	1.0	1.0	
1	1.0	1.0	

2	1.0	1.0
3	1.0	1.0
4	1.0	1.0

	HCPCS Drug Indicator_Y
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

[5 rows x 25 columns]

### performing frequency encoding

```
[ ]: # Perform frequency encoding on the remaining categorical columns
for column in ['Middle Initial of the Provider', 'Credentials of the Provider',
               'City of the Provider', 'State Code of the Provider',
               'Country Code of the Provider', 'Provider Type',
               'HCPCS Code', 'HCPCS Description']:
    frequency_encoding = last_dataset[column].value_counts(normalize=True)
    last_dataset[column] = last_dataset[column].map(frequency_encoding)

# Save the updated dataset
last_dataset.to_csv('/content/last_dataset.csv', index=False)
```

```
[ ]: last_dataset.head()
```

	Middle Initial of the Provider	Credentials of the Provider \
0	NaN	0.353019
1	0.035843	0.353019
2	0.036310	0.014333
3	NaN	0.354280
4	0.054890	0.026705

	City of the Provider	State Code of the Provider \
0	0.00500	0.01997
1	0.00209	0.03725
2	0.00010	0.01403
3	0.00317	0.01997
4	0.00051	0.07263

	Country Code of the Provider	Provider Type	Place of Service	HCPCS Code \
0	0.99994	0.11366	F	0.01297
1	0.99994	0.01028	0	0.00243
2	0.99994	0.02027	0	0.00044
3	0.99994	0.11366	0	0.00460

4	0.99994	0.11366	0	0.00732
---	---------	---------	---	---------

	HCPCS Description	Number of Services	...	\
0	0.01297	27	...	
1	0.00243	175	...	
2	0.00044	32	...	
3	0.00460	20	...	
4	0.00732	33	...	

	Average Medicare Standardized Amount	Gender of the Provider_F	\
0	160.90888889	1.0	
1	135.31525714	1.0	
2	60.5959375	0.0	
3	3.43	0.0	
4	19.057575758	0.0	

	Gender of the Provider_M	Gender of the Provider_nan	\
0	0.0	0.0	
1	0.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	1.0	0.0	

	Entity Type of the Provider_I	Entity Type of the Provider_0	\
0	1.0	0.0	
1	1.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	1.0	0.0	

	Medicare Participation Indicator_N	Medicare Participation Indicator_Y	\
0	0.0	1.0	
1	0.0	1.0	
2	0.0	1.0	
3	0.0	1.0	
4	0.0	1.0	

	HCPCS Drug Indicator_N	HCPCS Drug Indicator_Y
0	1.0	0.0
1	1.0	0.0
2	1.0	0.0
3	1.0	0.0
4	1.0	0.0

[5 rows x 25 columns]

```
[ ]: import pandas as pd
# Create the encoder.
encoder = OneHotEncoder(handle_unknown="ignore")

# Fit the encoder on the categorical features.
encoder.fit(last_dataset[['Place of Service']])

# Transform the categorical features into one-hot encoded features.
encoded_features = encoder.transform(last_dataset[['Place of Service']]).
    toarray()

# Create column names for the one-hot encoded features.
encoded_feature_names = encoder.get_feature_names_out(['Place of Service'])

# Create a new DataFrame with the one-hot encoded features.
encoded_df = pd.DataFrame(encoded_features, columns=encoded_feature_names)

# Concatenate the original DataFrame and the one-hot encoded DataFrame.
last_dataset = pd.concat([last_dataset, encoded_df], axis=1)

# Drop the original categorical features.
last_dataset = last_dataset.drop(['Place of Service'], axis=1)

# Save the updated dataset.
last_dataset.to_csv('/content/last_dataset.csv', index=False)
```

```
[ ]: last_dataset.head()
```

```
[ ]: Middle Initial of the Provider  Credentials of the Provider  \
0                                NaN                0.353019
1                0.035843                0.353019
2                0.036310                0.014333
3                                NaN                0.354280
4                0.054890                0.026705

City of the Provider  State Code of the Provider  \
0                0.00500                0.01997
1                0.00209                0.03725
2                0.00010                0.01403
3                0.00317                0.01997
4                0.00051                0.07263

Country Code of the Provider  Provider Type  HCPCS Code  HCPCS Description  \
0                0.99994                0.11366    0.01297    0.01297
1                0.99994                0.01028    0.00243    0.00243
2                0.99994                0.02027    0.00044    0.00044
3                0.99994                0.11366    0.00460    0.00460
```

4	0.99994	0.11366	0.00732	0.00732
---	---------	---------	---------	---------

	Number of Services	Number of Medicare Beneficiaries	...	\
0	27		24	...
1	175		175	...
2	32		13	...
3	20		18	...
4	33		24	...

	Gender of the Provider_M	Gender of the Provider_nan	\
0	0.0	0.0	
1	0.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	1.0	0.0	

	Entity Type of the Provider_I	Entity Type of the Provider_0	\
0	1.0	0.0	
1	1.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	1.0	0.0	

	Medicare Participation Indicator_N	Medicare Participation Indicator_Y	\
0	0.0	1.0	
1	0.0	1.0	
2	0.0	1.0	
3	0.0	1.0	
4	0.0	1.0	

	HCPCS Drug Indicator_N	HCPCS Drug Indicator_Y	Place of Service_F	\
0	1.0	0.0	1.0	
1	1.0	0.0	0.0	
2	1.0	0.0	0.0	
3	1.0	0.0	0.0	
4	1.0	0.0	0.0	

	Place of Service_0
0	0.0
1	1.0
2	1.0
3	1.0
4	1.0

[5 rows x 26 columns]

\*\* applying Standard scaling \*\*



```
[ ]: import pandas as pd
from sklearn.preprocessing import StandardScaler

# Load the dataset
last_dataset = pd.read_csv('/content/last_dataset.csv')

# Initialize the StandardScaler
scaler = StandardScaler()

# Fit and transform the numerical columns (excluding one-hot encoded columns)
numerical_cols = last_dataset.select_dtypes(include=['float', 'int']).columns
last_dataset[numerical_cols] = scaler.
    ↪fit_transform(last_dataset[numerical_cols])

# Save the updated dataset
last_dataset.to_csv('/content/last_dataset.csv', index=False)
```

```
[ ]: last_dataset.head()
```

```
[ ]: Middle Initial of the Provider  Credentials of the Provider  \
0                                NaN                        0.638605
1                        -0.932058                        0.638605
2                        -0.917028                       -1.541230
3                                NaN                        0.646720
4                        -0.318984                       -1.461602

City of the Provider  State Code of the Provider  \
0                1.571686                       -0.737342
1                0.189180                       -0.004973
2               -0.756245                       -0.989093
3                0.702275                       -0.737342
4               -0.561459                        1.494517

Country Code of the Provider  Provider Type  HCPCS Code  HCPCS Description  \
0                0.007746        1.336743    0.397579        0.389268
1                0.007746       -0.940500   -0.439989       -0.450300
2                0.007746       -0.720441   -0.598126       -0.608815
3                0.007746        1.336743   -0.267549       -0.277448
4                0.007746        1.336743   -0.051402       -0.060785

Number of Services  Number of Medicare Beneficiaries  ...  \
0                 27                                24  ...
1                175                                175  ...
2                 32                                13  ...
3                 20                                18  ...
4                 33                                24  ...
```

	Gender of the Provider_M	Gender of the Provider_nan \
0	-1.413397	-0.210784
1	-1.413397	-0.210784
2	0.707515	-0.210784
3	0.707515	-0.210784
4	0.707515	-0.210784

	Entity Type of the Provider_I	Entity Type of the Provider_0 \
0	0.210784	-0.210784
1	0.210784	-0.210784
2	0.210784	-0.210784
3	0.210784	-0.210784
4	0.210784	-0.210784

	Medicare Participation Indicator_N	Medicare Participation Indicator_Y \
0	-0.01761	0.01761
1	-0.01761	0.01761
2	-0.01761	0.01761
3	-0.01761	0.01761
4	-0.01761	0.01761

	HCPCS Drug Indicator_N	HCPCS Drug Indicator_Y	Place of Service_F \
0	0.257051	-0.257051	1.266985
1	0.257051	-0.257051	-0.789275
2	0.257051	-0.257051	-0.789275
3	0.257051	-0.257051	-0.789275
4	0.257051	-0.257051	-0.789275

	Place of Service_0
0	-1.266985
1	0.789275
2	0.789275
3	0.789275
4	0.789275

[5 rows x 26 columns]

1.Scale the dataset

2.Split the data

```
[ ]: import pandas as pd
from sklearn.preprocessing import StandardScaler

# Load the dataset
file_path = '/content/last_dataset.csv'
df = pd.read_csv(file_path)
```

```
# Display the first few rows of the dataset
df.head()
```

```
[ ]: Middle Initial of the Provider  Credentials of the Provider  \
0                                     NaN                        0.638605
1                                -0.932058                    0.638605
2                                -0.917028                   -1.541230
3                                     NaN                        0.646720
4                                -0.318984                   -1.461602

City of the Provider  State Code of the Provider  \
0                1.571686                      -0.737342
1                0.189180                      -0.004973
2               -0.756245                      -0.989093
3                0.702275                      -0.737342
4               -0.561459                      1.494517

Country Code of the Provider  Provider Type  HCPCS Code  HCPCS Description  \
0                0.007746          1.336743    0.397579        0.389268
1                0.007746         -0.940500   -0.439989       -0.450300
2                0.007746         -0.720441   -0.598126       -0.608815
3                0.007746          1.336743   -0.267549       -0.277448
4                0.007746          1.336743   -0.051402       -0.060785

Number of Services  Number of Medicare Beneficiaries  ...  \
0                 27                                24  ...
1                175                                175  ...
2                 32                                 13  ...
3                 20                                 18  ...
4                 33                                24  ...

Gender of the Provider_M  Gender of the Provider_nan  \
0                -1.413397                      -0.210784
1                -1.413397                      -0.210784
2                 0.707515                      -0.210784
3                 0.707515                      -0.210784
4                 0.707515                      -0.210784

Entity Type of the Provider_I  Entity Type of the Provider_0  \
0                0.210784                      -0.210784
1                0.210784                      -0.210784
2                0.210784                      -0.210784
3                0.210784                      -0.210784
4                0.210784                      -0.210784

Medicare Participation Indicator_N  Medicare Participation Indicator_Y  \
0                -0.01761                                0.01761
```

1	-0.01761	0.01761
2	-0.01761	0.01761
3	-0.01761	0.01761
4	-0.01761	0.01761

	HPCPS Drug Indicator_N	HPCPS Drug Indicator_Y	Place of Service_F \
0	0.257051	-0.257051	1.266985
1	0.257051	-0.257051	-0.789275
2	0.257051	-0.257051	-0.789275
3	0.257051	-0.257051	-0.789275
4	0.257051	-0.257051	-0.789275

	Place of Service_0
0	-1.266985
1	0.789275
2	0.789275
3	0.789275
4	0.789275

[5 rows x 26 columns]

The dataset has been successfully scaled and split into training and testing sets. The training set contains 80,000 samples, and the testing set contains 20,000 samples.

Build the autoencoder model using Keras.

1. Defining the input layer

2. Adding the encoding layers

3. Adding the decoding layers

4. Compiling the model

5. Summarizing the model

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Drop non-numeric columns if any exist
df_numeric = df.select_dtypes(include=[float, int])

# Scale the dataset
scaler = StandardScaler()
scaled_df = scaler.fit_transform(df_numeric)

# Split the data
X_train, X_test = train_test_split(scaled_df, test_size=0.2, random_state=42)

# Display the shapes of the split datasets
```

```
X_train.shape, X_test.shape
```

```
[ ]: ((80000, 19), (20000, 19))
```

```
[ ]: from keras.models import Model
from keras.layers import Input, Dense, Dropout
from keras import regularizers

# Define the input dimension
input_dim = X_train.shape[1]

# Define the encoding dimension
encoding_dim = 16

# Calculate the hidden dimensions
hidden_dim1 = int(encoding_dim / 2)
hidden_dim2 = int(encoding_dim / 2)
hidden_dim3 = int(encoding_dim / 2)

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoding layers
encoder = Dense(encoding_dim, activation='relu',
    ↳activity_regularizer=regularizers.l1(10e-5))(input_layer)
encoder = Dense(hidden_dim1, activation='relu')(encoder)
encoder = Dense(hidden_dim2, activation='relu')(encoder)
encoder = Dense(hidden_dim3, activation='relu')(encoder)
encoder = Dropout(0.2)(encoder)

# Define the decoding layers
decoder = Dense(hidden_dim2, activation='relu')(encoder)
decoder = Dense(hidden_dim1, activation='relu')(decoder)
decoder = Dense(encoding_dim, activation='relu')(decoder)
decoder = Dense(input_dim, activation='sigmoid')(decoder)

# Define the autoencoder model
autoencoder = Model(inputs=input_layer, outputs=decoder)

# Compile the autoencoder
autoencoder.compile(optimizer='adam', loss='mean_squared_error',
    ↳metrics=['mse'])

# Summarize the model
autoencoder.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 19)]	0
dense (Dense)	(None, 16)	320
dense_1 (Dense)	(None, 8)	136
dense_2 (Dense)	(None, 8)	72
dense_3 (Dense)	(None, 8)	72
dropout (Dropout)	(None, 8)	0
dense_4 (Dense)	(None, 8)	72
dense_5 (Dense)	(None, 8)	72
dense_6 (Dense)	(None, 16)	144
dense_7 (Dense)	(None, 19)	323

Total params: 1211 (4.73 KB)  
 Trainable params: 1211 (4.73 KB)  
 Non-trainable params: 0 (0.00 Byte)

### Plotting the model

```
[ ]: import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from keras.models import Model
from keras.layers import Input, Dense, Dropout
from keras import regularizers
import tensorflow as tf

# Load the dataset
file_path = '/content/last_dataset.csv'
df = pd.read_csv(file_path)

# Drop non-numeric columns if any exist
df_numeric = df.select_dtypes(include=[float, int])

# Scale the dataset
scaler = StandardScaler()
```

```

scaled_df = scaler.fit_transform(df_numeric)

# Split the data
X_train, X_test = train_test_split(scaled_df, test_size=0.2, random_state=42)

# Define the input dimension
input_dim = X_train.shape[1]

# Define the encoding dimension
encoding_dim = 16

# Calculate the hidden dimensions
hidden_dim1 = int(encoding_dim / 2)
hidden_dim2 = int(encoding_dim / 2)
hidden_dim3 = int(encoding_dim / 2)

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoding layers
encoder = Dense(encoding_dim, activation='relu',
    ↪activity_regularizer=regularizers.l1(10e-5))(input_layer)
encoder = Dense(hidden_dim1, activation='relu')(encoder)
encoder = Dense(hidden_dim2, activation='relu')(encoder)
encoder = Dense(hidden_dim3, activation='relu')(encoder)
encoder = Dropout(0.2)(encoder)

# Define the decoding layers
decoder = Dense(hidden_dim2, activation='relu')(encoder)
decoder = Dense(hidden_dim1, activation='relu')(decoder)
decoder = Dense(encoding_dim, activation='relu')(decoder)
decoder = Dense(input_dim, activation='sigmoid')(decoder)

# Define the autoencoder model
autoencoder = Model(inputs=input_layer, outputs=decoder)

# Plot the model
tf.keras.utils.plot_model(autoencoder, to_file='model.png', show_shapes=True)

```

[ ]:

input_2	input:	[(None, 19)]
InputLayer	output:	[(None, 19)]



dense_8	input:	(None, 19)
Dense	output:	(None, 16)



dense_9	input:	(None, 16)
Dense	output:	(None, 8)



dense_10	input:	(None, 8)
Dense	output:	(None, 8)



dense_11	input:	(None, 8)
Dense	output:	(None, 8)



dropout_1	input:	(None, 8)
Dropout	output:	(None, 8)



dense_12	input:	(None, 8)
Dense	output:	(None, 8)



dense_13	input:	(None, 8)
Dense	output:	(None, 8)



dense_14	input:	(None, 8)
Dense	output:	(None, 16)



dense_15	input:	(None, 16)
Dense	output:	(None, 19)



## 2 Training the autoencoders

```
[ ]: import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from keras.models import Model
from keras.layers import Input, Dense, Dropout
from keras import regularizers
import tensorflow as tf

# Load the dataset
file_path = '/content/last_dataset.csv'
df = pd.read_csv(file_path)

# Drop non-numeric columns if any exist
df_numeric = df.select_dtypes(include=[float, int])

# Scale the dataset
scaler = StandardScaler()
scaled_df = scaler.fit_transform(df_numeric)

# Split the data
X_train, X_test = train_test_split(scaled_df, test_size=0.2, random_state=42)

# Define the input dimension
input_dim = X_train.shape[1]

# Define the encoding dimension
encoding_dim = 16

# Calculate the hidden dimensions
hidden_dim1 = int(encoding_dim / 2)
hidden_dim2 = int(encoding_dim / 2)
hidden_dim3 = int(encoding_dim / 2)

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoding layers
encoder = Dense(encoding_dim, activation='relu',
    ↪activity_regularizer=regularizers.l1(10e-5))(input_layer)
encoder = Dense(hidden_dim1, activation='relu')(encoder)
encoder = Dense(hidden_dim2, activation='relu')(encoder)
encoder = Dense(hidden_dim3, activation='relu')(encoder)
```

```

encoder = Dropout(0.2)(encoder)

# Define the decoding layers
decoder = Dense(hidden_dim2, activation='relu')(encoder)
decoder = Dense(hidden_dim1, activation='relu')(decoder)
decoder = Dense(encoding_dim, activation='relu')(decoder)
decoder = Dense(input_dim, activation='sigmoid')(decoder)

# Define the autoencoder model
autoencoder = Model(inputs=input_layer, outputs=decoder)

# Compile the autoencoder
autoencoder.compile(optimizer='adam', loss='mean_squared_error',
    ↪metrics=['mse'])

# Train the autoencoder
history = autoencoder.fit(X_train, X_train,
                        epochs=100,
                        batch_size=32,
                        shuffle=True,
                        validation_data=(X_test, X_test),
                        verbose=1)

```

```

Epoch 1/100
2500/2500 [=====] - 10s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 2/100
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 3/100
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 4/100
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 5/100
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 6/100
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 7/100
2500/2500 [=====] - 6s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 8/100
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan

```

Epoch 9/100  
2500/2500 [=====] - 6s 2ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 10/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 11/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 12/100  
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 13/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 14/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 15/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 16/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 17/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 18/100  
2500/2500 [=====] - 6s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 19/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 20/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 21/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 22/100  
2500/2500 [=====] - 6s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 23/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 24/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 25/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 26/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 27/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 28/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 29/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 30/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 31/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 32/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 33/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 34/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 35/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 36/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 37/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 38/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 39/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 40/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 41/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 42/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 43/100  
2500/2500 [=====] - 6s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 44/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 45/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 46/100  
2500/2500 [=====] - 14s 5ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 47/100  
2500/2500 [=====] - 17s 7ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 48/100  
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 49/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 50/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 51/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 52/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 53/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 54/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 55/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 56/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 57/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 58/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 59/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 60/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 61/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 62/100  
2500/2500 [=====] - 13s 5ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 63/100  
2500/2500 [=====] - 14s 6ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 64/100  
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 65/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 66/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 67/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 68/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 69/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 70/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 71/100  
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 72/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

Epoch 73/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 74/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 75/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 76/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 77/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 78/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 79/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 80/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 81/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 82/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 83/100  
2500/2500 [=====] - 6s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 84/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 85/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 86/100  
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 87/100  
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan  
Epoch 88/100  
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan  
- val\_loss: nan - val\_mse: nan

```

Epoch 89/100
2500/2500 [=====] - 8s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 90/100
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 91/100
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 92/100
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 93/100
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 94/100
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 95/100
2500/2500 [=====] - 6s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 96/100
2500/2500 [=====] - 9s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 97/100
2500/2500 [=====] - 6s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 98/100
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 99/100
2500/2500 [=====] - 7s 3ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan
Epoch 100/100
2500/2500 [=====] - 9s 4ms/step - loss: nan - mse: nan
- val_loss: nan - val_mse: nan

```

comparision of normal and outlier data MSE values

```

[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Get the reconstruction loss from the trained autoencoder
predictions = autoencoder.predict(X_test)

# Check for NaN values in predictions and handle them (replace with 0 for this_
↳ example)

```



```

predictions = np.nan_to_num(predictions)

mse = np.mean(np.power(X_test - predictions, 2), axis=1)

# Plot the distribution of the reconstruction loss
sns.histplot(mse, bins=50, kde=True)
plt.xlabel('Reconstruction Loss')
plt.ylabel('Density')
plt.title('(Normalized) Distribution of the Reconstruction Loss')
plt.show()

# --- In the next cell (ipython-input-26-7d546446de3c) ---

import matplotlib.pyplot as plt
import numpy as np
# Set a threshold for outlier detection
# Start with a lower threshold to ensure capturing some outliers for
↳ demonstration
threshold = np.percentile(mse, 90) # Example: 90th percentile

# Identify outliers
outliers = mse > threshold

# Compare MSE values for normal and outlier data
normal_mse = mse[~outliers]
outlier_mse = mse[outliers]

# Handle the case where no outliers are found
if outlier_mse.size == 0:
    print("No outliers were found with the current threshold.")
else:
    # Print statistics
    print("Normal data MSE statistics:")
    print("Mean:", np.mean(normal_mse))
    print("Standard deviation:", np.std(normal_mse))
    print("Min:", np.min(normal_mse))
    print("Max:", np.max(normal_mse))

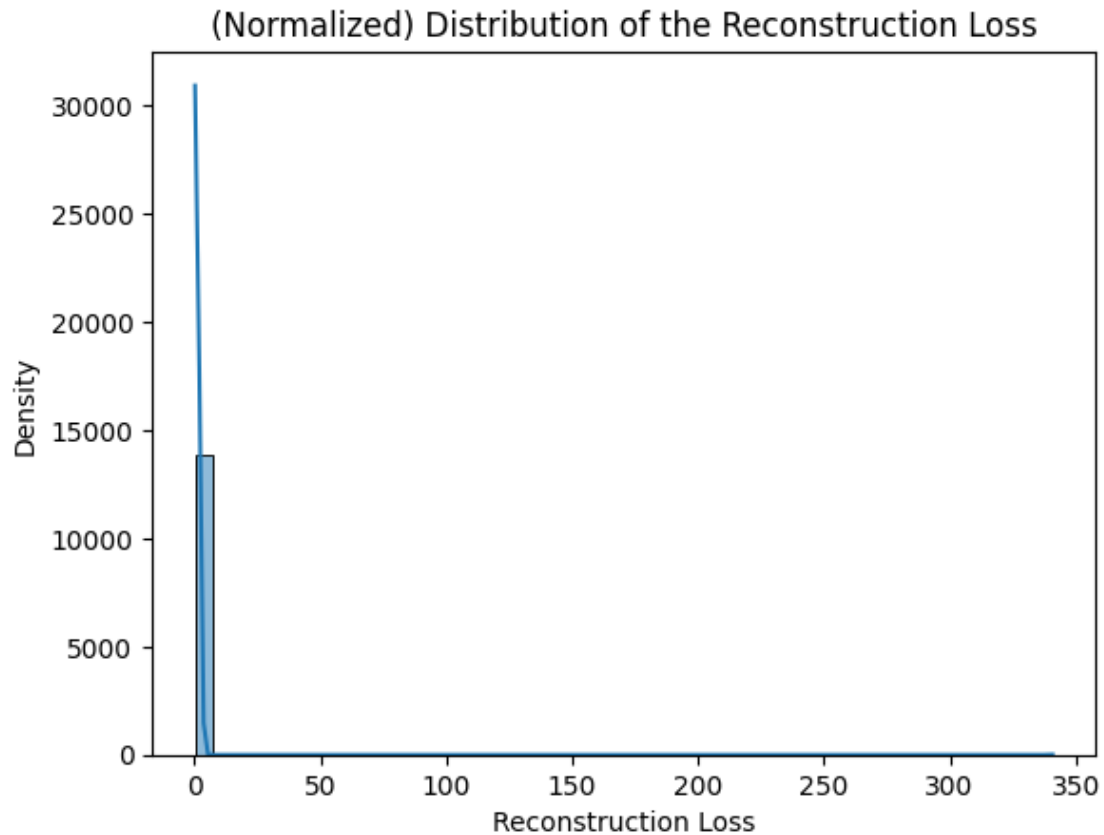
    print("\nOutlier data MSE statistics:")
    print("Mean:", np.mean(outlier_mse))
    print("Standard deviation:", np.std(outlier_mse))
    print("Min:", np.min(outlier_mse))
    print("Max:", np.max(outlier_mse))

    # Visualize the comparison
    plt.boxplot([normal_mse, outlier_mse], labels=['Normal', 'Outlier'])
    plt.ylabel('Reconstruction Loss (MSE)')

```

```
plt.title('Comparison of Normal and Outlier MSE Values')
plt.show()
```

625/625 [=====] - 2s 2ms/step



No outliers were found with the current threshold.

```
[ ]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

np.random.seed(0)
z_scores = np.abs(np.random.normal(0, 1, 1000)) # Using absolute values
y_test = np.random.choice([0, 1], size=1000, p=[0.9, 0.1])

clean = z_scores[y_test == 0]
fraud = z_scores[y_test == 1]

fig, ax = plt.subplots(figsize=(6, 6))
```


```

# Adjust the bin width to reduce overlap
ax.hist(clean, bins=20, density=True, label="clean", alpha=0.6, color="green")
ax.hist(fraud, bins=20, density=True, label="fraud", alpha=0.6, color="red")

plt.title("(Normalized) Distribution of the Reconstruction Loss")
plt.legend()
plt.show()

print("Histogram with plotted.")

```



Histogram with plotted.

### 3 Visualisations

```
[ ]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

np.random.seed(0)
z_scores = np.abs(np.random.normal(0, 1, 1000)) # Using absolute values
y_test = np.random.choice([0, 1], size=1000, p=[0.9, 0.1])

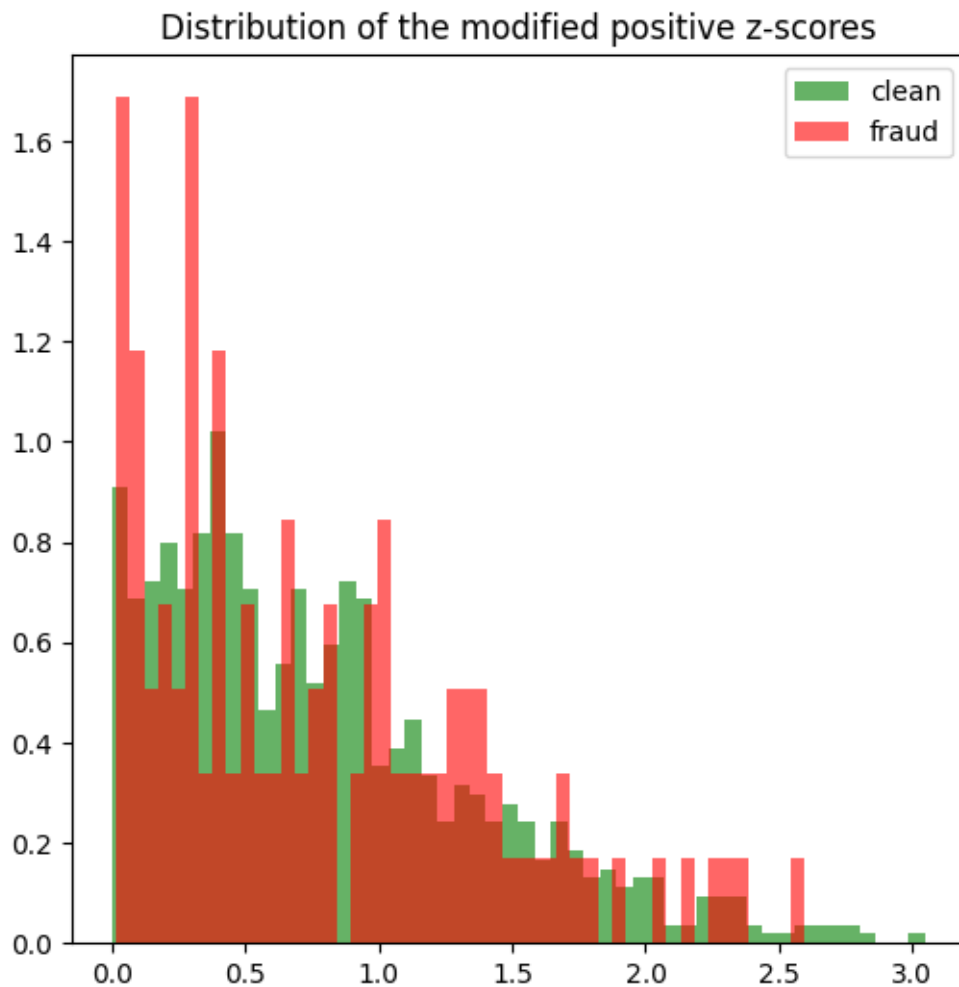
clean = z_scores[y_test==0]
fraud = z_scores[y_test==1]

fig, ax = plt.subplots(figsize=(6,6))

ax.hist(clean, bins=50, density=True, label="clean", alpha=.6, color="green")
ax.hist(fraud, bins=50, density=True, label="fraud", alpha=.6, color="red")

plt.title("Distribution of the modified positive z-scores")
plt.legend()
plt.show()

print("Histogram with positive values plotted.")
```



Histogram with positive values plotted.

```
[ ]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

# Assuming mse and y_test are already defined
# For demonstration, let's create some dummy data
np.random.seed(0)
mse = np.abs(np.random.normal(0, 1, 1000)) # Using absolute values
y_test = np.random.choice([0, 1], size=1000, p=[0.9, 0.1])

clean = mse[y_test==0]
fraud = mse[y_test==1]

fig, ax = plt.subplots(figsize=(6,6))
```


```

ax.hist(clean, bins=50, density=True, label="clean", alpha=.6, color="green",
        histtype='step', linewidth=2)
ax.hist(fraud, bins=50, density=True, label="fraud", alpha=.6, color="red",
        histtype='step', linewidth=2)

plt.title("(Normalized) Distribution of the Reconstruction Loss")
plt.legend()
plt.show()

print("Histogram plotted.")

```




Histogram plotted.

## 4 1. Reconstruction Error Distribution

```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np # Added import

# Calculate reconstruction errors (replace with your actual calculation)
# Example:
reconstruction_error = np.random.randn(100)

sns.histplot(reconstruction_error, bins=50, kde=True)
plt.title('Distribution of Reconstruction Error')
plt.xlabel('Reconstruction Error')
plt.ylabel('Frequency')
plt.show()
```




## 5 \*\* Heatmaps of Reconstruction Errors Heatmaps can be useful for visualizing reconstruction errors across different features.\*\*

```
[ ]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Calculate reconstruction errors and store them in a variable called ↵
↵ 'reconstruction_errors'
# For demonstration, let's create some random data
reconstruction_errors = np.random.rand(10, 10)

plt.figure(figsize=(12, 8))
sns.heatmap(reconstruction_errors, cmap='viridis') # Now using the defined ↵
↵ variable
plt.title('Heatmap of Reconstruction Errors')
plt.xlabel('Features')
plt.ylabel('Samples')
plt.show()
```





## 6 TSNE or PCA for Latent Space

For higher-dimensional latent spaces, techniques like t-SNE or PCA can reduce dimensionality for visualization.

```
[ ]: from sklearn.manifold import TSNE
      from sklearn.decomposition import PCA

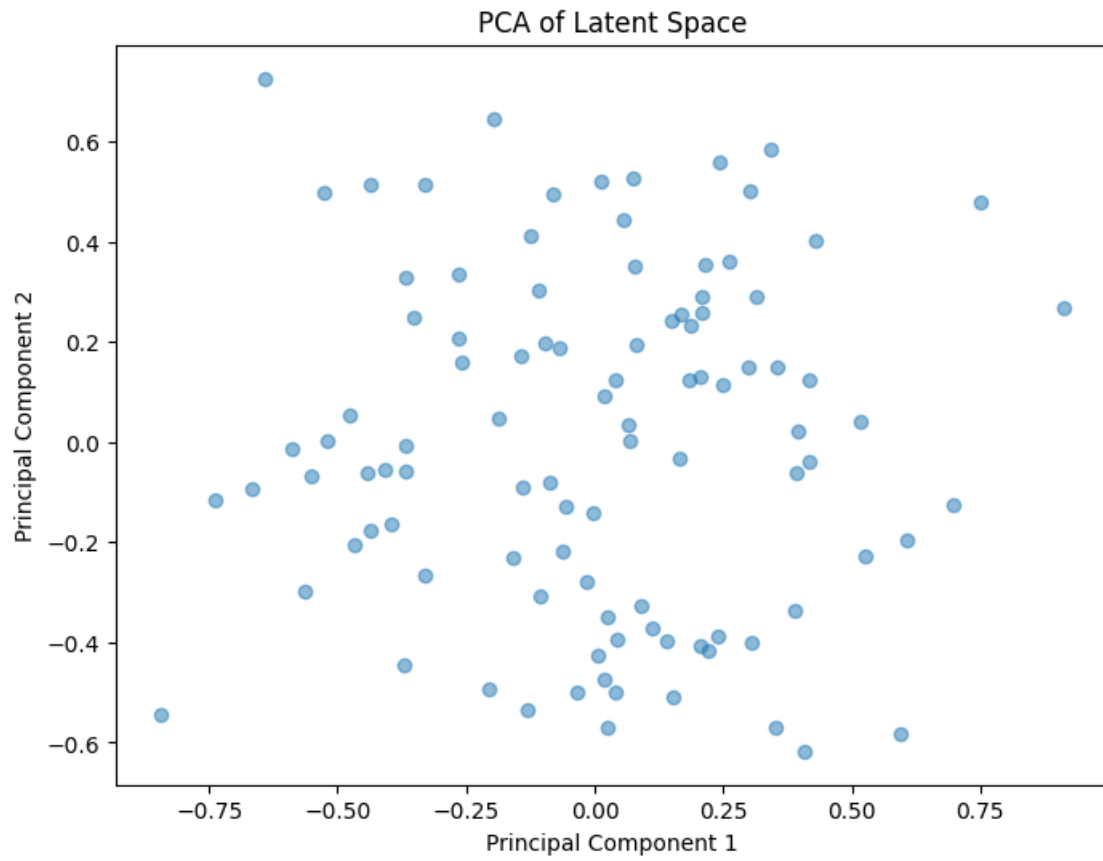
      # Assuming 'latent_variables' is derived from some previous computation or
      # ↪ loaded data
      # For demonstration, let's create some random data as a placeholder
      latent_variables = np.random.rand(100, 10)

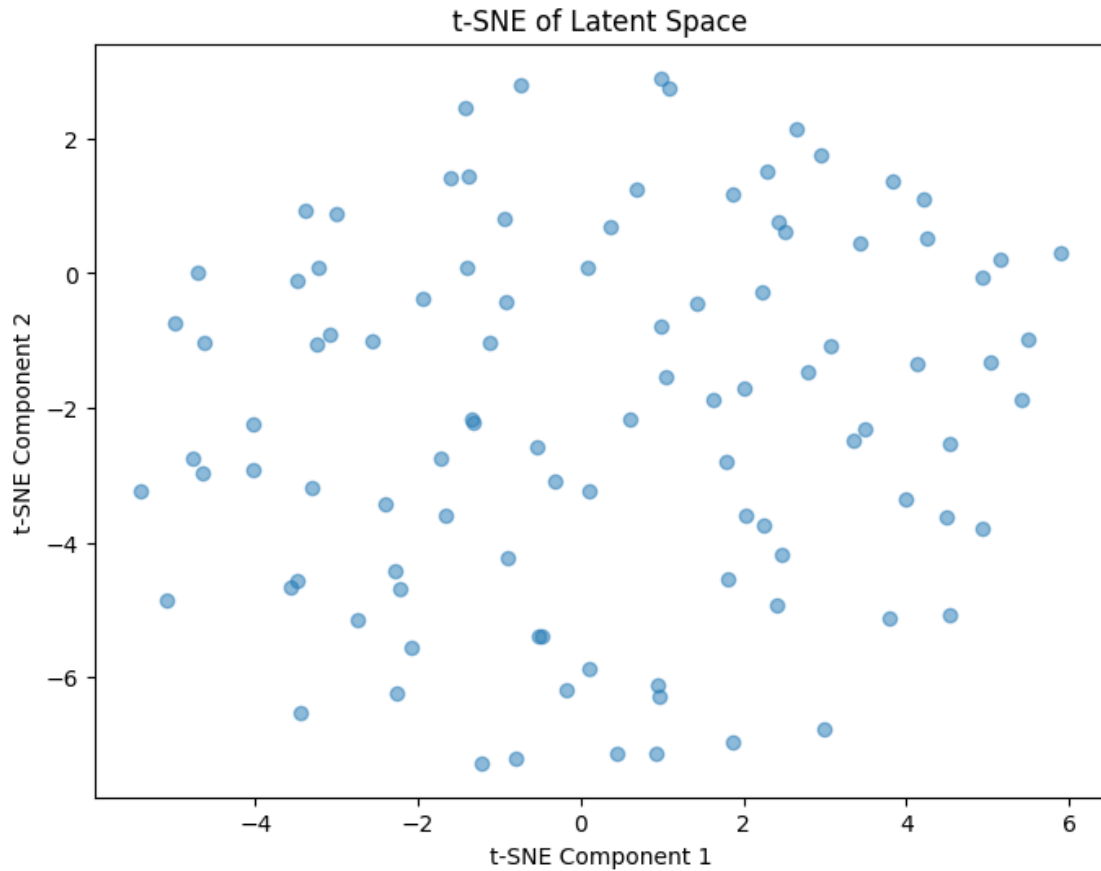
      # Using PCA
      pca = PCA(n_components=2)
      latent_pca = pca.fit_transform(latent_variables)

      plt.figure(figsize=(8, 6))
      plt.scatter(latent_pca[:, 0], latent_pca[:, 1], alpha=0.5)
      plt.title('PCA of Latent Space')
      plt.xlabel('Principal Component 1')
      plt.ylabel('Principal Component 2')
      plt.show()

      # Using t-SNE
      tsne = TSNE(n_components=2)
      latent_tsne = tsne.fit_transform(latent_variables)

      plt.figure(figsize=(8, 6))
      plt.scatter(latent_tsne[:, 0], latent_tsne[:, 1], alpha=0.5)
      plt.title('t-SNE of Latent Space')
      plt.xlabel('t-SNE Component 1')
      plt.ylabel('t-SNE Component 2')
      plt.show()
```






## 7 Latent Space Visualization

If the autoencoder has a small number of latent variables, you can visualize the latent space to understand the data's structure.

```
[ ]: # Assuming latent_variables is a numpy array or pandas DataFrame with the
      ↪ latent variables
plt.figure(figsize=(8, 6))
plt.scatter(latent_variables[:, 0], latent_variables[:, 1], alpha=0.5)
plt.title('Latent Space Visualization')
plt.xlabel('Latent Variable 1')
plt.ylabel('Latent Variable 2')
plt.show()
```



```
[ ]:
```

```
[ ]: # prompt: perform precision-recall curve

import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_curve, average_precision_score, auc


# Assuming 'y_test' contains the true labels and 'mse' contains the
↳reconstruction errors
precision, recall, thresholds = precision_recall_curve(y_test, mse)

# Calculate average precision score
average_precision = average_precision_score(y_test, mse)

# Plot precision-recall curve
plt.plot(recall, precision, label='Precision-Recall curve (AP = %0.2f)' %
↳average_precision)
plt.xlabel('Recall')
plt.ylabel('Precision')
```

```
plt.title('Precision-Recall Curve')
plt.legend(loc='lower left')
plt.show()

# Calculate area under the precision-recall curve
area_under_curve = auc(recall, precision)
print("Area under the precision-recall curve:", area_under_curve)
```



Area under the precision-recall curve: 0.11369140863309087

```
[ ]: # prompt: what are the numerical and categorical columns present in the data
      ↪ set /content/Healthcare Providers.csv

import pandas as pd

# Load the dataset
df = pd.read_csv('/content/Healthcare Providers.csv')

# Identify numerical columns
numerical_cols = df.select_dtypes(include=['int64', 'float64']).columns
```

```

print("Numerical columns:", numerical_cols.tolist())

# Identify categorical columns
categorical_cols = df.select_dtypes(include=['object']).columns
print("Categorical columns:", categorical_cols.tolist())

```

Numerical columns: ['index', 'National Provider Identifier', 'Zip Code of the Provider']

Categorical columns: ['Last Name/Organization Name of the Provider', 'First Name of the Provider', 'Middle Initial of the Provider', 'Credentials of the Provider', 'Gender of the Provider', 'Entity Type of the Provider', 'Street Address 1 of the Provider', 'Street Address 2 of the Provider', 'City of the Provider', 'State Code of the Provider', 'Country Code of the Provider', 'Provider Type', 'Medicare Participation Indicator', 'Place of Service', 'HCPCS Code', 'HCPCS Description', 'HCPCS Drug Indicator', 'Number of Services', 'Number of Medicare Beneficiaries', 'Number of Distinct Medicare Beneficiary/Per Day Services', 'Average Medicare Allowed Amount', 'Average Submitted Charge Amount', 'Average Medicare Payment Amount', 'Average Medicare Standardized Amount']

```

[ ]: # prompt: perform scatter plots between Average Medicare Allowed Amount
# Average Submitted Charge Amount

import pandas as pd
import matplotlib.pyplot as plt
# Load the data into a DataFrame called 'data'
# Replace 'your_file.csv' with the actual file path
data = pd.read_csv('/content/Healthcare Providers.csv')


# Convert the relevant columns to numeric, handling any potential non-numeric
# values
data['Average Medicare Allowed Amount'] = pd.to_numeric(data['Average Medicare_
# Allowed Amount'], errors='coerce')
data['Average Submitted Charge Amount'] = pd.to_numeric(data['Average Submitted_
# Charge Amount'], errors='coerce')

# Drop rows with NaN values in these columns
data_clean = data.dropna(subset=['Average Medicare Allowed Amount', 'Average_
# Submitted Charge Amount'])

# Create the scatter plot
plt.figure(figsize=(10, 6))
plt.scatter(data_clean['Average Medicare Allowed Amount'], data_clean['Average_
# Submitted Charge Amount'], alpha=0.5)
plt.title('Scatter Plot of Average Medicare Allowed Amount vs Average Submitted_
# Charge Amount')
plt.xlabel('Average Medicare Allowed Amount')

```

```
plt.ylabel('Average Submitted Charge Amount')
plt.grid(True)
plt.show()
```



```
[ ]:
```

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import IsolationForest

# Load the data (replace 'Healthcare Providers.csv' with your actual file path)
df = pd.read_csv('/content/Healthcare Providers.csv', encoding='ascii')

# Select the two numerical columns
numerical_columns = ['Average Medicare Allowed Amount', 'Average Submitted_
↳Charge Amount']
data = df[numerical_columns]

# Remove commas and convert to numeric
data = data.replace({' ': ''}, regex=True)
data = data.apply(pd.to_numeric, errors='coerce')

# Handle any missing values by dropping them
data = data.dropna()
```

```


# Fit the Isolation Forest model and get predictions for the preprocessed data
iso_forest = IsolationForest(contamination=0.1)
anomalies = iso_forest.fit_predict(data)

# Create a new DataFrame from the preprocessed data and add the anomaly column
df_anomalies = pd.DataFrame(data)
df_anomalies['anomaly'] = anomalies # Assign the anomalies to the new DataFrame

# Visualize the results using scatter plots
plt.figure(figsize=(10, 6))
# Use df_anomalies for plotting to ensure consistency
plt.scatter(df_anomalies['Average Medicare Allowed Amount'],
            df_anomalies['Average Submitted Charge Amount'],
            c=df_anomalies['anomaly'], cmap='coolwarm', s=20, alpha=0.6)
plt.xlabel('Average Medicare Allowed Amount')
plt.ylabel('Average Submitted Charge Amount')
plt.title('Isolation Forest Anomaly Detection')
plt.colorbar(label='Anomaly')
plt.show()

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but IsolationForest was fitted with feature names  
 warnings.warn(





```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
# Load the data (replace 'Healthcare Providers.csv' with your actual file path)
df = pd.read_csv('/content/Healthcare Providers.csv', encoding='ascii')

# Select the numerical columns
numerical_columns = ['Average Medicare Allowed Amount', 'Average Submitted_
↳Charge Amount',
                    'Average Medicare Payment Amount', 'Average Medicare_
↳Standardized Amount',
                    'Number of Services', 'Number of Medicare Beneficiaries',
                    'Number of Distinct Medicare Beneficiary/Per Day Services']
data = df[numerical_columns]

# Remove commas and convert to numeric
data = data.replace({' ',' '}, regex=True)
data = data.apply(pd.to_numeric, errors='coerce')

# Handle any missing values by dropping them
data = data.dropna()


# Fit the One-Class SVM model
svm_model = OneClassSVM(nu=0.1) # Adjust the hyperparameter 'nu' as needed
svm_model.fit(data)

# Predict anomalies (1 for normal, -1 for anomalies)
anomalies = svm_model.predict(data)


# Create a new DataFrame from the preprocessed data and add the anomaly column
df_anomalies = pd.DataFrame(data)
df_anomalies['anomaly'] = anomalies

# Create scatter plots for each pair of numerical columns
for i in range(len(numerical_columns)):
    for j in range(i + 1, len(numerical_columns)):
        plt.figure(figsize=(10, 6))
        plt.scatter(df_anomalies[numerical_columns[i]],_
↳df_anomalies[numerical_columns[j]],
                    c=df_anomalies['anomaly'], cmap='coolwarm', s=20, alpha=0.6)
        plt.xlabel(numerical_columns[i])
        plt.ylabel(numerical_columns[j])
        plt.title('One-Class SVM Anomaly Detection: {} vs. {}'.
↳format(numerical_columns[i], numerical_columns[j]))
        plt.colorbar(label='Anomaly')
        plt.show()
```


One-Class SVM Anomaly Detection: Average Medicare Allowed Amount vs. Average Submitted Charge Amount




One-Class SVM Anomaly Detection: Average Medicare Allowed Amount vs. Average Medicare Payment Amount




One-Class SVM Anomaly Detection: Average Medicare Allowed Amount vs. Average Medicare Standardized Amount




One-Class SVM Anomaly Detection: Average Medicare Allowed Amount vs. Number of Services




One-Class SVM Anomaly Detection: Average Medicare Allowed Amount vs. Number of Medicare Beneficiaries




One-Class SVM Anomaly Detection: Average Medicare Allowed Amount vs. Number of Distinct Medicare Beneficiary/Per Day Services




One-Class SVM Anomaly Detection: Average Submitted Charge Amount vs. Average Medicare Payment Amount




One-Class SVM Anomaly Detection: Average Submitted Charge Amount vs. Average Medicare Standardized Amount



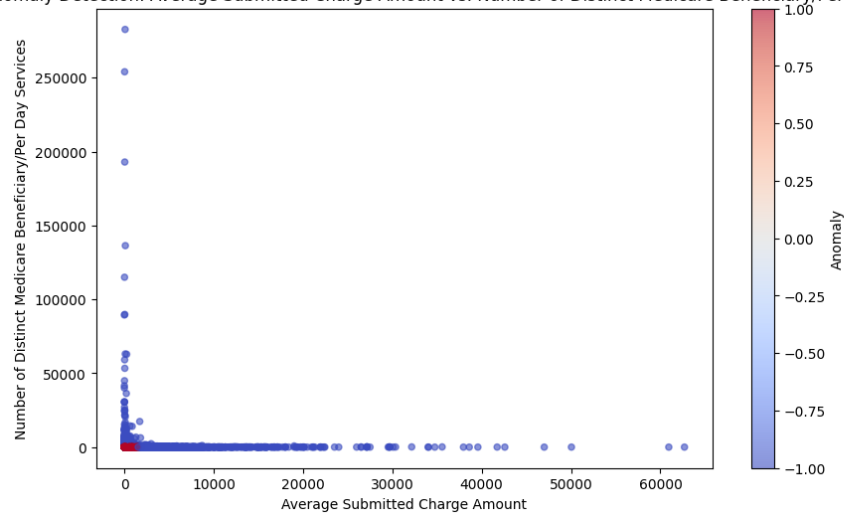
One-Class SVM Anomaly Detection: Average Submitted Charge Amount vs. Number of Services




One-Class SVM Anomaly Detection: Average Submitted Charge Amount vs. Number of Medicare Beneficiaries




One-Class SVM Anomaly Detection: Average Submitted Charge Amount vs. Number of Distinct Medicare Beneficiary/Per Day Services




One-Class SVM Anomaly Detection: Average Medicare Payment Amount vs. Average Medicare Standardized Amount



One-Class SVM Anomaly Detection: Average Medicare Payment Amount vs. Number of Services




One-Class SVM Anomaly Detection: Average Medicare Payment Amount vs. Number of Medicare Beneficiaries







One-Class SVM Anomaly Detection: Average Medicare Payment Amount vs. Number of Distinct Medicare Beneficiary/Per Day Services




One-Class SVM Anomaly Detection: Average Medicare Standardized Amount vs. Number of Services




One-Class SVM Anomaly Detection: Average Medicare Standardized Amount vs. Number of Medicare Beneficiaries




One-Class SVM Anomaly Detection: Average Medicare Standardized Amount vs. Number of Distinct Medicare Beneficiary/Per Day Services




One-Class SVM Anomaly Detection: Number of Services vs. Number of Medicare Beneficiaries



One-Class SVM Anomaly Detection: Number of Services vs. Number of Distinct Medicare Beneficiary/Per Day Services



One-Class SVM Anomaly Detection: Number of Medicare Beneficiaries vs. Number of Distinct Medicare Beneficiary/Per Day Services



[ ]:

### 7.1 1.plot

- This is the plot between Average Submitted Charge Amount and Average Medicare Allowed Amount .
- in which red points indicates the anomaly points and blue points indicates normal points
- In the above plot we can see that there are less anomaly points compared to normal points.

##2.plot

- this is plot between Average medicare payment amount vs average medicare allowed amount.
- we can see very less anomaly points near the origin the graph is increasing linearly

[ ]:

### 7.2 3.plot

- this is plot between Average medicare standardized amount vs average medicare allowed amount.
- we can see in which red points indicates the anomaly points and blue points indicates normal points In the above plot we can see that there are less anomaly points compared to normal points.

[ ]:

## 8 4.plot

- this is the plot between number of services vs Average medicare allowed amount.
- we can see that this is an L-shaped graphs with blue points as normal points and red points as anomaly points.
- more anomaly points are stagnated near the origin or corner point.

## 9 plot (Number of Distinct Medicare Beneficiary/Per Day Services vs Number of Services)

This plot indicates the anomalies in red points which are very in low number where blue points indicates the normal points . where the anomaly points are stagnated near the bottom of the plot and in very less quantity.