

## MileStone2

### Name- Om Late

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
data = pd.read_csv("/content/Health1.csv")
data.head()
```



	index	National Provider Identifier	Last Name/Organization Name of the Provider	First Name of the Provider	Middle Initial of the Provider	Credentials of the Provider	Gende of th Provide
0	8774979	1891106191	UPADHYAYULA	SATYASREE	NaN	M.D.	
1	3354385	1346202256	JONES	WENDY	P	M.D.	
2	3001884	1306820956	DUROCHER	RICHARD	W	DPM	I
3	7594822	1770523540	FULLARD	JASPER	NaN	MD	I
4	746159	1073627758	PERROTTI	ANTHONY	E	DO	I

5 rows × 27 columns

```
# original data
data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 67325 entries, 0 to 67324
Data columns (total 27 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   index                                     67325 non-null  int64
1   National Provider Identifier              67325 non-null  int64
```

2	Last Name/Organization Name of the Provider	67325	non-null	object
3	First Name of the Provider	64477	non-null	object
4	Middle Initial of the Provider	47580	non-null	object
5	Credentials of the Provider	62477	non-null	object
6	Gender of the Provider	64478	non-null	object
7	Entity Type of the Provider	67325	non-null	object
8	Street Address 1 of the Provider	67325	non-null	object
9	Street Address 2 of the Provider	27507	non-null	object
10	City of the Provider	67325	non-null	object
11	Zip Code of the Provider	67325	non-null	float64
12	State Code of the Provider	67324	non-null	object
13	Country Code of the Provider	67324	non-null	object
14	Provider Type	67324	non-null	object
15	Medicare Participation Indicator	67324	non-null	object
16	Place of Service	67324	non-null	object
17	HCPCS Code	67324	non-null	object
18	HCPCS Description	67324	non-null	object
19	HCPCS Drug Indicator	67324	non-null	object
20	Number of Services	67324	non-null	object
21	Number of Medicare Beneficiaries	67324	non-null	object
22	Number of Distinct Medicare Beneficiary/Per Day Services	67324	non-null	object
23	Average Medicare Allowed Amount	67324	non-null	object
24	Average Submitted Charge Amount	67324	non-null	object
25	Average Medicare Payment Amount	67324	non-null	object
26	Average Medicare Standardized Amount	67324	non-null	object

dtypes: float64(1), int64(2), object(24)  
memory usage: 13.9+ MB

```

irrelevant_columns=['Entity Type of the Provider',
                    'Street Address 1 of the Provider',
                    'Street Address 2 of the Provider',
                    'Zip Code of the Provider',
                    'Medicare Participation Indicator',
                    'Place of Service',
                    'HCPCS Code',
                    'HCPCS Description',
                    'HCPCS Drug Indicator',
                    'Country Code of the Provider']
data=data.drop(columns=irrelevant_columns)

```

```
data.head()
```



	index	National Provider Identifier	Last Name/Organization Name of the Provider	First Name of the Provider	Middle Initial of the Provider	Credentials of the Provider	Gender of the Provider
0	8774979	1891106191	UPADHYAYULA	SATYASREE	NaN	M.D.	
1	3354385	1346202256	JONES	WENDY	P	M.D.	
2	3001884	1306820956	DUROCHER	RICHARD	W	DPM	F
3	7594822	1770523540	FULLARD	JASPER	NaN	MD	F
4	746159	1073627758	PERROTTI	ANTHONY	E	DO	F

```
# Merging the name columns into a single column
data['Full Name'] = data['First Name of the Provider'].fillna('') + ' ' + \
    data['Middle Initial of the Provider'].fillna('') + ' ' + \
    data['Last Name/Organization Name of the Provider'].fillna('')
data['Full Name'] = data['Full Name'].str.strip()
data = data.drop(columns=['Last Name/Organization Name of the Provider',
    'First Name of the Provider',
    'Middle Initial of the Provider'])
full_name_column = data.pop('Full Name')
data.insert(1, 'Full Name', full_name_column)
data.head()
```



	index	Full Name	National Provider Identifier	Credentials of the Provider	Gender of the Provider	City of the Provider	State Code of the Provider
0	8774979	SATYASREE UPADHYAYULA	1891106191	M.D.	F	SAINT LOUIS	MO
1	3354385	WENDY P JONES	1346202256	M.D.	F	FAYETTEVILLE	NC
2	3001884	RICHARD W DUROCHER	1306820956	DPM	M	NORTH HAVEN	CT
3	7594822	JASPER FULLARD	1770523540	MD	M	KANSAS CITY	MO
4	746159	ANTHONY E PERROTTI	1073627758	DO	M	JUPITER	FL

```
# Uniform format of credentials
data['Credentials of the Provider'] = data['Credentials of the Provider'].str.replace(r'\s', '')
data.head()
```



	index	Full Name	National Provider Identifier	Credentials of the Provider	Gender of the Provider	City of the Provider	State Code of the Provider
0	8774979	SATYASREE UPADHYAYULA	1891106191	MD	F	SAINT LOUIS	MO
1	3354385	WENDY P JONES	1346202256	MD	F	FAYETTEVILLE	NC
2	3001884	RICHARD W DUROCHER	1306820956	DPM	M	NORTH HAVEN	CT
3	7594822	JASPER FULLARD	1770523540	MD	M	KANSAS CITY	MO
4	746159	ANTHONY E PERROTTI	1073627758	DO	M	JUPITER	FL

## ✓ Converting Object to Numeric

```
numeric_columns = [
    'Number of Services',
    'Number of Medicare Beneficiaries',
    'Number of Distinct Medicare Beneficiary/Per Day Services',
    'Average Medicare Allowed Amount',
    'Average Submitted Charge Amount',
    'Average Medicare Payment Amount',
    'Average Medicare Standardized Amount'
]
for column in numeric_columns:
    data[column] = pd.to_numeric(data[column], errors='coerce')
```

```
data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 67325 entries, 0 to 67324
Data columns (total 15 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   index                                     67325 non-null  int64
1   Full Name                                67325 non-null  object
2   National Provider Identifier              67325 non-null  int64
3   Credentials of the Provider               62477 non-null  object
4   Gender of the Provider                    64478 non-null  object
```

5	City of the Provider	67325	non-null	object
6	State Code of the Provider	67324	non-null	object
7	Provider Type	67324	non-null	object
8	Number of Services	65519	non-null	float6
9	Number of Medicare Beneficiaries	67039	non-null	float6
10	Number of Distinct Medicare Beneficiary/Per Day Services	66316	non-null	float6
11	Average Medicare Allowed Amount	66821	non-null	float6
12	Average Submitted Charge Amount	62847	non-null	float6
13	Average Medicare Payment Amount	67004	non-null	float6
14	Average Medicare Standardized Amount	67006	non-null	float6

dtypes: float64(7), int64(2), object(6)  
memory usage: 7.7+ MB

```
# missing values
print(data.isnull().sum())
```

```
➡ index                                0
Full Name                              0
National Provider Identifier             0
Credentials of the Provider             4848
Gender of the Provider                  2847
City of the Provider                    0
State Code of the Provider              1
Provider Type                           1
Number of Services                      1806
Number of Medicare Beneficiaries        286
Number of Distinct Medicare Beneficiary/Per Day Services 1009
Average Medicare Allowed Amount         504
Average Submitted Charge Amount         4478
Average Medicare Payment Amount         321
Average Medicare Standardized Amount    319
dtype: int64
```

```
# Imputation of numeric missing values with mean
data[numeric_columns] = data[numeric_columns].fillna(data[numeric_columns].mean())
print(data.isnull().sum())
```

```
➡ index                                0
Full Name                              0
National Provider Identifier             0
Credentials of the Provider             4848
Gender of the Provider                  2847
City of the Provider                    0
State Code of the Provider              1
Provider Type                           1
Number of Services                      0
Number of Medicare Beneficiaries        0
Number of Distinct Medicare Beneficiary/Per Day Services 0
Average Medicare Allowed Amount         0
Average Submitted Charge Amount         0
Average Medicare Payment Amount         0
Average Medicare Standardized Amount    0
dtype: int64
```

```
categorical_columns = ['Credentials of the Provider',
                        'Gender of the Provider',
                        'City of the Provider',
                        'State Code of the Provider']

for column in categorical_columns:
    data[column].fillna(data[column].mode()[0], inplace=True)

print(data.isnull().sum())
```

index	0
Full Name	0
National Provider Identifier	0
Credentials of the Provider	0
Gender of the Provider	0
City of the Provider	0
State Code of the Provider	0
Provider Type	1
Number of Services	0
Number of Medicare Beneficiaries	0
Number of Distinct Medicare Beneficiary/Per Day Services	0
Average Medicare Allowed Amount	0
Average Submitted Charge Amount	0
Average Medicare Payment Amount	0
Average Medicare Standardized Amount	0
dtype: int64	

```
# Check for duplicates
print(data.duplicated().sum())
```

0
---

```
data.head()
```

	index	Full Name	National Provider Identifier	Credentials of the Provider	Gender of the Provider	City of the Provider	State Code of the Provider
0	8774979	SATYASREE UPADHYAYULA	1891106191	MD	F	SAINT LOUIS	MO
1	3354385	WENDY P JONES	1346202256	MD	F	FAYETTEVILLE	NC
2	3001884	RICHARD W DUROCHER	1306820956	DPM	M	NORTH HAVEN	CT
3	7594822	JASPER FULLARD	1770523540	MD	M	KANSAS CITY	MO
4	746159	ANTHONY E PERROTTI	1073627758	DO	M	JUPITER	FL

```
def frequency_encode(df, columns):
    for column in columns:
        freq_encoding = df[column].value_counts() / len(df)
        new_column_name = column + '_Freq'
        # Check if column exists before inserting
        if new_column_name not in df.columns:
            df.insert(df.columns.get_loc(column) + 1, new_column_name, df[column].map(freq_en
    return df
columns_to_encode=[ 'Credentials of the Provider',
                    'Gender of the Provider',
                    'Provider Type',
                    'State Code of the Provider']
data = frequency_encode(data, columns_to_encode)
data.head()
```



	index	Full Name	National Provider Identifier	Credentials of the Provider	Credentials of the Provider_Freq	Gender of the Provider	Gender o Provider.
0	8774979	SATYASREE UPADHYAYULA	1891106191	MD	0.735130	F	0.7
1	3354385	WENDY P JONES	1346202256	MD	0.735130	F	0.7
2	3001884	RICHARD W DUROCHER	1306820956	DPM	0.020160	M	0.7
3	7594822	JASPER FULLARD	1770523540	MD	0.735130	M	0.7
4	746159	ANTHONY E PERROTTI	1073627758	DO	0.064157	M	0.7

```

#Performing Standardization on Numerical Columns
from sklearn.preprocessing import StandardScaler
standardization_columns=['Number of Services',
                          'Number of Medicare Beneficiaries',
                          'Number of Distinct Medicare Beneficiary/Per Day Services',
                          'Average Medicare Allowed Amount',
                          'Average Submitted Charge Amount',
                          'Average Medicare Payment Amount',
                          'Average Medicare Standardized Amount',
                          'Credentials of the Provider_Freq',
                          'Gender of the Provider_Freq',
                          'State Code of the Provider_Freq' ]

# Standardization
standard_scaler = StandardScaler()
data[standardization_columns] = standard_scaler.fit_transform(data[standardization_columns])
data_copy=data.copy()
print("Standardized DataFrame:")
data.head()

```


➡ Standardized DataFrame:

	index	Full Name	National Provider Identifier	Credentials of the Provider	Credentials of the Provider_Freq	Gender of the Provider	Gender o f Provider.
0	8774979	SATYASREE UPADHYAYULA	1891106191	MD	0.599764	F	-1.51
1	3354385	WENDY P JONES	1346202256	MD	0.599764	F	-1.51
2	3001884	RICHARD W DUROCHER	1306820956	DPM	-1.669524	M	0.61
3	7594822	JASPER FULLARD	1770523540	MD	0.599764	M	0.61
4	746159	ANTHONY E PERROTTI	1073627758	DO	-1.529881	M	0.61




```
#Dimensionality Reduction using PCA
from sklearn.decomposition import PCA
df=data.copy()
# Imputation of categorical columns with mode
categorical_columns = ['Full Name',
                       'Credentials of the Provider',
                       'Gender of the Provider',
                       'City of the Provider',
                       'Provider Type',
                       'State Code of the Provider']


for column in df.columns:
    df[column].fillna(df[column].mode()[0], inplace=True)
df = df.drop(columns=categorical_columns)
pca = PCA(n_components=2)
pca_result = pca.fit_transform(df)
# DataFrame of PCA results
pca_df = pd.DataFrame(pca_result, columns=['PCA1', 'PCA2'])
# Scatter plot of PCA1 and PCA2
plt.figure(figsize=(8, 6))
plt.scatter(pca_df['PCA1'], pca_df['PCA2'])
plt.xlabel('PCA1')
plt.ylabel('PCA2')
plt.title('PCA of Transformed Data')
plt.grid(True)
plt.show()
```




```
# Plot PCA1 as a histogram
plt.hist(pca_df['PCA1'], bins=20, edgecolor='black')
plt.xlabel('PCA1 Values')
plt.ylabel('Frequency')
plt.title('Histogram of PCA1 Values')
plt.grid(True)
plt.show()
```




```
# Plot PCA2 as a histogram
plt.hist(pca_df['PCA2'], bins=20, edgecolor='black')
plt.xlabel('PCA2 Values')
plt.ylabel('Frequency')
plt.title('Histogram of PCA2 Values')
plt.grid(True)
plt.show()
```



```
#CLUSTERING
#K MEANS CLUSTERING
from sklearn.cluster import KMeans, DBSCAN
from sklearn.metrics import silhouette_score
kmeans = KMeans(n_clusters=2, random_state=42)
data['Cluster_KMeans'] = kmeans.fit_predict(data[numeric_columns])
sns.scatterplot(data=data, x='Number of Medicare Beneficiaries', y='Average Submitted Cha
    hue='Cluster_KMeans', palette='viridis', legend='full')
plt.title('K-Means Clustering')
plt.show()
```

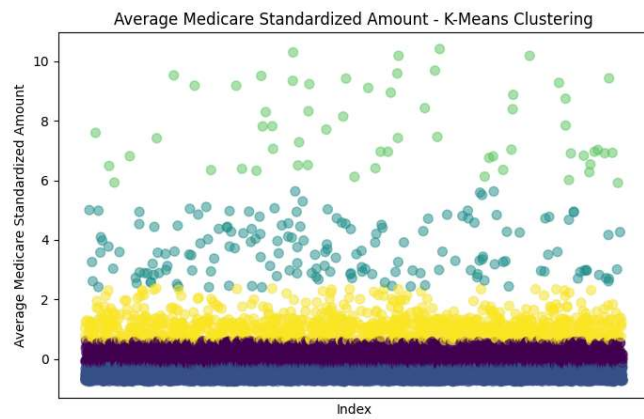
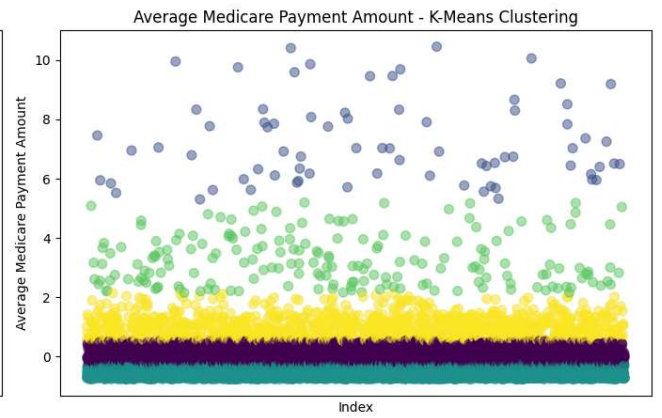



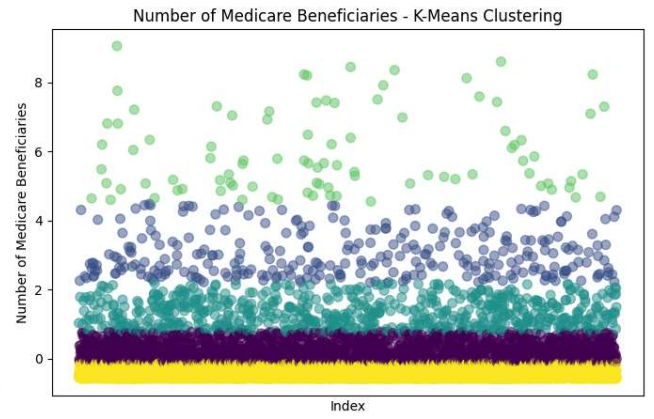
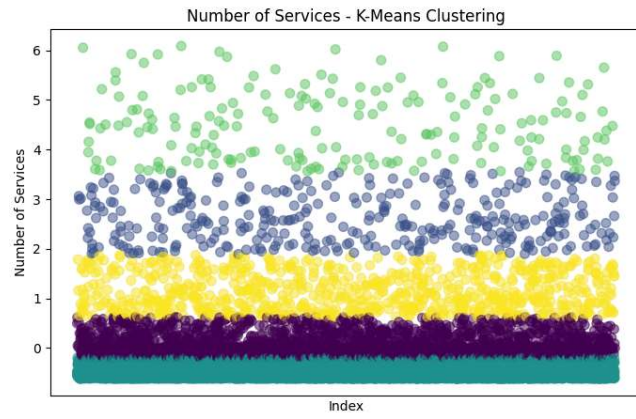


```
# Clustering using K-Means
kmeans = KMeans(n_clusters=3, random_state=42)
data['Cluster_KMeans'] = kmeans.fit_predict(data[numeric_columns])
sns.scatterplot(data=data, x='Number of Services', y='Average Medicare Payment Amount', h
    palette='viridis', legend='full')
plt.title('K-Means Clustering')
plt.show()
```



```
#Algoplot of K-Means
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
k = 5
fig, axes = plt.subplots(4, 2, figsize=(14, 18))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
axes = axes.flatten()
for i, col in enumerate(numeric_columns):
    # Perform K-Means clustering on the current column
    kmeans = KMeans(n_clusters=k, random_state=0)
    data['Cluster'] = kmeans.fit_predict(data[[col]])
    # Plot the column against its K-Means cluster assignments
    ax = axes[i]
    ax.scatter(data.index, data[col], c=data['Cluster'], s=50, alpha=0.5)
    ax.set_title(f'{col} - K-Means Clustering')
    ax.set_xlabel('Index')
    ax.set_ylabel(col)
    ax.set_xticks([])
if i < len(numeric_columns) - 2:
    ax.set_xticklabels([])
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```



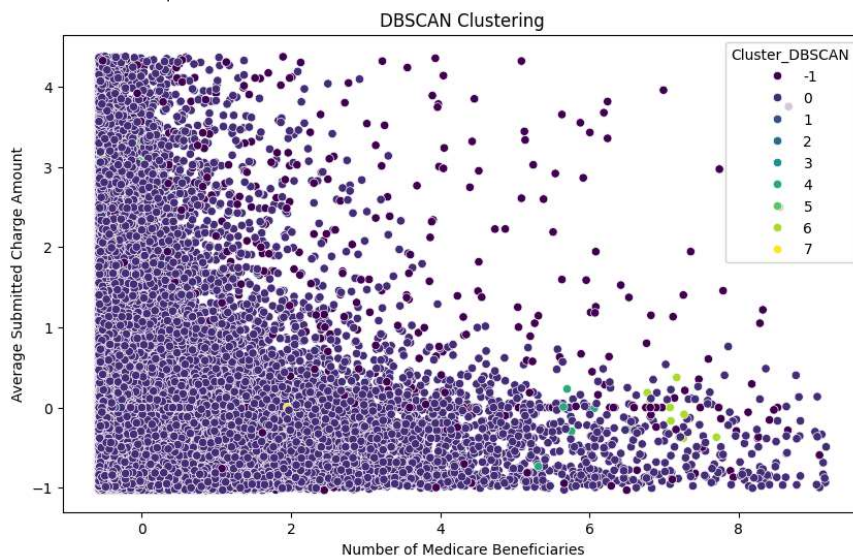


```
#DB SCAN CLUSTERING
from sklearn.cluster import DBSCAN
# Clustering using DBSCAN
dbscan = DBSCAN(eps=0.7, min_samples=6)
data['Cluster_DBSCAN'] = dbscan.fit_predict(data[numeric_columns])
# Number of noise points
num_noise_points = (data['Cluster_DBSCAN'] == -1).sum()
print(f"Number of noise points: {num_noise_points}")
plt.figure(figsize=(10, 6))
sns.scatterplot(data=data, x='Number of Medicare Beneficiaries', y='Average Submitted Cha
    hue='Cluster_DBSCAN', palette='viridis', legend='full')
plt.title('DBSCAN Clustering')
plt.show()
```



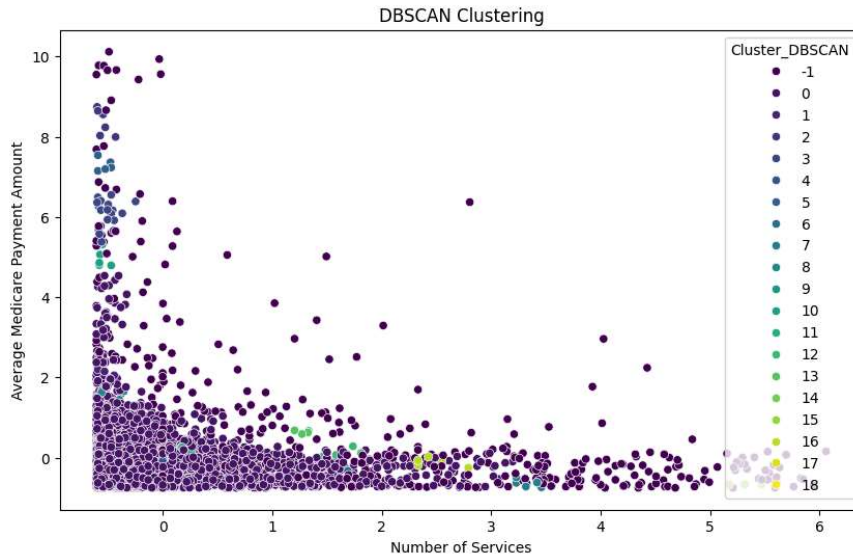
```
#DB SCAN CLUSTERING
from sklearn.cluster import DBSCAN
# Clustering using DBSCAN
dbscan = DBSCAN(eps=0.7, min_samples=6)
data['Cluster_DBSCAN'] = dbscan.fit_predict(data[numeric_columns])
# Number of noise points
num_noise_points = (data['Cluster_DBSCAN'] == -1).sum()
print(f"Number of noise points: {num_noise_points}")
plt.figure(figsize=(10, 6))
sns.scatterplot(data=data, x='Number of Medicare Beneficiaries', y='Average Submitted Charge Amount',
                hue='Cluster_DBSCAN', palette='viridis', legend='full')
plt.title('DBSCAN Clustering')
plt.show()
```

→ Number of noise points: 703



```
dbscan = DBSCAN(eps=0.5, min_samples=4)
data['Cluster_DBSCAN'] = dbscan.fit_predict(data[numeric_columns])
num_noise_points = (data['Cluster_DBSCAN'] == -1).sum()
print(f"Number of noise points: {num_noise_points}")
plt.figure(figsize=(10, 6))
sns.scatterplot(data=data, x='Number of Services', y='Average Medicare Payment Amount',
                hue='Cluster_DBSCAN', palette='viridis', legend='full')
plt.title('DBSCAN Clustering')
plt.show()
```

➡ Number of noise points: 445



```
#Algoplot of DBScan
eps = 0.5
min_samples = 3
data = data[numeric_columns].dropna()
data = data.sample(n=5000, random_state=42)
fig, axes = plt.subplots(4, 2, figsize=(14, 18))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
axes = axes.flatten()
for i, col in enumerate(numeric_columns):
    # Perform DBSCAN clustering on the current column
    dbscan = DBSCAN(eps=eps, min_samples=min_samples)
    # Reshape data to 2D array for DBSCAN
    data_col = data[[col]].values.reshape(-1, 1)
    data['Cluster'] = dbscan.fit_predict(data_col)
    # Plot the column against its DBSCAN cluster assignments
    ax = axes[i]
    scatter = ax.scatter(data.index, data[col], c=data['Cluster'], cmap='viridis', s=50, alpha=0.5)
    ax.set_title(f'{col} - DBSCAN Clustering')

    cbar = plt.colorbar(scatter, ax=ax)
    cbar.set_label('Cluster')
if i < len(numeric_columns) - 2:
    ax.set_xticklabels([])
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

